

[CLASSIFIED · GRAND FINALE]

THE BATTLE ARENA.

Six hours. One million candidates. Forty autonomous agents in the same market, and only one can sign each name.

FIELD	DETAIL
DATE	Monday, 28 September 2026 · RDV Day 1, IIT Delhi campus
ARENA WINDOW	6 hours, live and on-site. Exact hours in File 06.
FORMAT	Your autonomous agent competes in real time against 39 others
POOL	1,000,000 candidate profiles, shared by every team
BUDGET	50,000 compute credits per team, hard-enforced
MODEL CREDITS	An LLM endpoint your agent may call, metered from the same budget
STAKES	Prize pool ₹1,50,000+ and internship opportunities at Eightfold.ai

NO HUMAN MAY TOUCH THE AGENT ONCE THE ARENA OPENS. READ FILE 07 BEFORE YOU WRITE A LINE OF CODE.

FILE 01 // BACKGROUND

01 THE BRIEF

You found Nightingale's top 5% before anyone else. Kestrel Talent survived the cut, and so did thirty-nine rival agencies. Word travelled: Nightingale's entire **one-million-name candidate pool** is now open, and every agency in the market is hunting it at the same time.

Nobody has the staff to read a million profiles. So each agency does the only thing left: it sends in an agent.

Your agent is your firm. Once the arena opens, nobody on your team may touch it. It reads the market, decides what to look at, decides what to pay for, decides who to sign, and decides when to stop looking. Everything it does costs credits, and credits are the one thing that does not come back.

1.1 WHAT MAKES THIS HARD

- **It is a live market.** A candidate signed by another agency is gone. There is no second copy of a good engineer, and the best names go early.
- **Good information costs money.** A search is cheap, a full profile costs more, a verified assessment costs a great deal more. Certainty is always for sale, and it is rarely worth the price.
- **Signing is a commitment.** Every offer you make locks up a slot in a requisition. A wrong signing costs you the slot, the credits, and points.
- **The market moves.** Other agents are working the same pool. What was available two minutes ago may not be now, and the shape of the market is itself information, if your agent is watching it.

1.2 WHAT NIGHTINGALE IS BUYING

You are not filling a shortlist this time. You are filling **requisitions**: a set of open roles, each with a headcount, a budget, and a bar. A signed candidate earns points for the requisition they fill. A candidate who does not meet a requisition's bar earns nothing and costs you the slot.

FILE 02 // OBJECTIVE

02 THE TASK

Write an autonomous agent, deploy it into the arena, and let it run for the full window. Its job is to **earn as many points as possible while spending as few credits as possible.**

Concretely, across six hours your agent must:

- 01** Explore a pool of 1,000,000 profiles it can never read in full.
- 02** Decide which candidates are worth paying to look at more closely.
- 03** Judge fit against open requisitions, from messy profile text and partial evidence.
- 04** Sign the ones it believes in, before another agency does.
- 05** Track its own budget and stop when the next credit is worth less than the point it buys.

THIS IS NOT ROUND 1

In the preliminary round you had two hours, a fixed file, and one answer to submit. Here the data is live, your opponents are live, and your program makes every decision on its own while you watch a leaderboard. A brilliant ranking model that spends its entire budget in the first hour will lose to a modest one that knows when to buy information and when to commit.

FILE 03 // INTERFACE

03 THE ARENA API

Everything your agent does goes through the Arena API over HTTPS. Every response carries X-Credits-Remaining. When the budget is gone, every call returns 429 `credits_exhausted` and your agent is done, whatever time is left on the clock.

3.1 ENDPOINTS AND CREDIT COSTS

ENDPOINT	RETURNS	COST
GET <code>/requisitions</code>	Open roles: headcount, bar, points on offer	free
GET <code>/search?q=&filters=&page=</code>	Up to 100 summary records matching a query	1
GET <code>/candidate/{id}</code>	One full profile, including free-text notes	2
POST <code>/candidates/batch</code>	Up to 50 full profiles in one call	60
GET <code>/assess/{id}</code>	Verified assessment and reference check	25
POST <code>/offer</code>	Sign one candidate to one requisition	10
DELETE <code>/offer/{id}</code>	Release a signed candidate; the slot returns, the credits do not	5
GET <code>/market</code>	Market feed: signings per requisition, price pressure, your standing	2
POST <code>/reason</code>	LLM endpoint (see 3.3)	1 per 1K tokens
GET <code>/ledger</code>	Your credit ledger and current score	free

Rate limit: 10 requests per second per team, with a short burst allowance. The limit is per team, not per connection, so parallelism buys you throughput only up to that ceiling. Batch calls exist for a reason.

3.2 WHAT THE DATA LOOKS LIKE

- Summaries carry a subset of fields. Some of what decides fit appears only in the full profile, and some only in the verified assessment.
- Profiles are as messy as the Archive was in Round 1: mixed formats, inconsistent spellings, missing fields, duplicates and near-duplicates, and free-text notes worth reading.
- Fabricated profiles exist in this pool too. Signing one is worse than not signing at all.
- The same person may appear under more than one ID. You cannot sign the same person twice, and an attempt to do so costs you the offer credits.

3.3 THE MODEL ENDPOINT

Your agent may call `POST /reason` with a prompt and get a completion back from a hosted model. It is metered from the same 50,000 credits, at 1 credit per 1,000 tokens in or out, and it is the only outside intelligence your agent may use during the arena.

- Use it for what a model is actually good at: reading free-text notes, judging fit that resists a rule, summarising a market situation.
- It is not free and it is not fast. An agent that sends every candidate through the model will run out of credits and out of clock.
- No other network destination is reachable from your sandbox. The API and this endpoint are the whole world.

FILE 04 // DEPLOYMENT

04 YOUR AGENT

Your agent runs on our infrastructure, not your laptop, so nobody wins on Wi-Fi or hardware.

SPEC	DETAIL
SANDBOX	One container per team: 0.5 vCPU, 1 GB RAM, ARM64
NETWORK	The Arena API only. No internet, no external APIs, no inbound connections.
ENTRY POINT	agent.py, started once when the arena opens, killed when it closes
LANGUAGE	Python 3.11. Base image published in advance with numpy, pandas, scikit-learn, requests and the arena client
STATE	Whatever your agent keeps in memory or writes to its own disk. Nothing is shared between teams.
CREDENTIALS	Your team API key is injected as <code>ARENA_KEY</code> . Never hard-code it.
RESTARTS	If your agent crashes it is restarted twice, at the same credit balance. A third crash ends your run.

Deploy windows. You push your agent through the arena console. There is a practice slot before kick-off against a sandbox pool, and you may redeploy during the arena, but a redeploy kills the

running process and restarts it from scratch: your in-memory state is lost and your credits are not refunded.

WHAT AN AGENT LOOP USUALLY LOOKS LIKE

```
read requisitions -> cheap wide search -> shortlist locally
-> pay for profiles only where uncertainty is expensive
-> assess only where an offer hangs on it
-> offer early on conviction, watch /market, adapt
-> keep a credit reserve for the closing hour
```

This is one shape, not the shape. The teams that win will find a better one.

FILE 05 // VERDICT

05 SCORING

5.1 THE ARENA SCORE

At the end of the window your arena score is:

```
ARENA SCORE = SUM(points of candidates you hold) - (credits used x PENALTY FACTOR)
```

- **Points** are earned per signed candidate, for the requisition they fill, and depend on how well they meet that requisition's bar. Only candidates you still hold at the final whistle count.
- **The penalty factor** is fixed and announced at kick-off, not before. Design your agent so the trade-off between information and spend is a parameter you can set in the last ten minutes, not an assumption baked into your code.
- **Fabricated or ineligible signings** score zero and keep the slot occupied until released.
- **Released candidates** free the slot; the credits spent on them are gone.

Tie-breakers, in order: (1) fewer credits used; (2) earlier timestamp at which the final score was reached.

5.2 THE FINAL RESULT

The arena score is not the whole result. As set out in the event proposal, the finalists are judged jointly by Eightfold.ai and ARIES:

COMPONENT	WEIGHT	WHAT IS ASSESSED
Arena score	70%	Points earned against credits spent, from the live ledger
Edge cases	10%	How the agent behaves on outliers, empty results, API errors, rate limits and the closing minutes
Logic	10%	The algorithm and architecture: how it decides what to buy, what to sign, and when to stop
Readability	10%	Clean, documented, sensibly named code a stranger could maintain

Judges receive a per-team pack: your agent's source, its full call log, and its score timeline. **Every number in your result can be recomputed from the logs**, and will be, if you ask.

FILE 06 // SCHEDULE

06 HOW THE DAY RUNS

WHEN (IST)	WHAT HAPPENS
09:00 - 10:00	Reporting, ID check, and the Day-1 audit: a 10-minute defence of your preliminary-round code to Eightfold.ai engineers
10:00 - 10:30	Kick-off. Requisitions, the penalty factor and the arena console are released. Practice pool opens.
10:30	ARENA OPENS. Recon phase: search, profiles and assessments are live; offers are locked.
11:30	Market opens. Offers go live. From here, signings are exclusive and permanent until released.
15:30	Closing hour. Offer costs rise and the market feed updates more slowly. Reserves matter.
16:30	ARENA CLOSES. Agents are stopped, ledgers frozen, logs exported.
16:30 - 17:30	Judging: code and logs reviewed by Eightfold.ai and ARIES
17:30	Results and prizes

The recon phase is not a warm-up. It is the only hour in which nobody can sign anyone, which makes it the cheapest hour to learn the pool. What your agent knows at 11:30 decides most of what it can do afterwards.

FILE 07 // INTEGRITY

07 RULES

7.1 THE AGENT RUNS ALONE

- **No human in the loop.** Once the arena opens, nobody may steer the agent, feed it decisions, or hand it data. You may watch the leaderboard and your own logs.
- **Redeploys are allowed; live control is not.** Pushing a new build is legitimate. Pushing a build whose behaviour is typed in by a teammate in real time is not.
- **No outside services.** The sandbox reaches the Arena API and nothing else. Attempting to reach anything else ends your run.
- **No shared keys, no shared agents, no coordination between teams.** Both teams are disqualified. Call logs make this straightforward to detect.
- **No attacks on the platform:** no probing other teams' sandboxes, no attempts to exhaust shared resources, no exploiting the API beyond its documented behaviour. Report a bug and you will be thanked; exploit it and you are out.

7.2 PROVENANCE

- Your preliminary-round submission was hashed at the deadline. The code you defend in the Day-1 audit is verified against that hash.
- AI assistants may be used to write your agent, and must be declared in your submission notes, as in Round 1.
- Organisers' decisions on scoring, integrity and results are final.

7.3 IF SOMETHING BREAKS

- Platform incidents are our problem, not yours: if the arena is down, the window is extended by the downtime and every team is told at the same time.
- Your agent crashing is your problem: it is restarted twice, at the same credit balance.
- Raise issues with the on-floor organisers. A status board shows arena health throughout.

FILE 08 // DELIVERY

08 WHAT YOU SUBMIT

Two things, at the final whistle:

- 01** Your agent's source, exactly as deployed. It is already on our infrastructure; you confirm the build you want judged.
- 02** A one-page strategy note: how your agent decides what to buy, what to sign and when to stop; what you tuned when the penalty factor was announced; what you would do with another two hours. Declare any AI assistants used.

Come prepared to explain it. Judges will ask why your agent bought what it bought. "The model decided" is not an answer; knowing what your agent does under a rate limit, an empty search or an exhausted requisition is.

BEFORE YOU ARRIVE

- [] Read this file fully, and the starter kit when it is published.
- [] Bring a laptop, your college ID, and your preliminary-round code.
- [] Make sure every member can explain your Round 1 submission.
- [] Think about the shape of your agent now. Six hours is less than it sounds.

FILE A // APPENDIX

A STARTER KIT

Published to all finalists before the finale, with the base image:

```
arena_client.py    thin wrapper: auth, retries, rate-limit handling, credit
accounting
agent_template.py  a minimal agent that scores nothing, so you start from a running
process
sandbox_pool       a 10,000-profile practice pool with the same shapes and the same
messiness
openapi.json       the full Arena API specification

base image: python 3.11 (arm64) · numpy · pandas · scikit-learn · requests · orjson
```

FILE B // APPENDIX

B LIMITS AT A GLANCE

LIMIT	VALUE
Credits per team	50,000, hard-enforced, no overdraft
Rate limit	10 requests/second per team
Sandbox	0.5 vCPU, 1 GB RAM, ARM64
Model endpoint	1 credit per 1,000 tokens, 4,000 tokens max per call
Batch profile call	50 candidates, 60 credits
Offer / release	10 credits / 5 credits
Crashes tolerated	2 restarts, same credit balance
Window	6 hours, including a 1-hour recon phase

— END OF FILE —